Technical Report:
Nonparametric Partial Dependence

Terence Parr · James D. Wilson

Abstract Partial dependence curves introduced by [3] (FPD), are an important model interpretation tool, but are not accessible to business analysts and scientists without the skills to choose, tune, and assess machine learning models. It is also common for the same partial dependence algorithm on the same data to give meaningfully different curves for different models, which calls into question their validity. Expertise is required to distinguish between model artifacts and true relationships in the data.

In this paper, we contribute methods for computing partial dependence curves, for both numerical (StratPD) and categorical variables (CatStratPD), that work directly from training data rather than predictions of a model. We provide evidence of biased results with FPD, ALE ([1]), and SHAP ([11]), plus demonstrate that our approach works correctly on synthetic and real data sets. Our goal is not to argue that model-based techniques are not useful. Rather, we hope to open a new line of inquiry into nonparametric partial dependence.

1 Introduction

Partial dependence, the isolated effect of a specific variable or variables on the response variable, y, is important to researchers and practitioners in many disparate fields such as medicine, business, and the social sciences. For example, in medicine, researchers are interested in the relationship between an individual's demographics or clinical features and their susceptibility to illness. Business analysts at a car manufacturer might need to know how changes in their supply chain are affecting defect rates. Climate scientists are interested in how different atmospheric carbon levels affect temperature.

For an explanatory matrix, $\mathbf{X}$, with a single variable, x_1, a plot of the y against x_1 visualizes the marginal effect of feature x_1 on y exactly. Given two or more features, one can similarly plot the marginal effects of each feature separately, however, the analysis is complicated by the interactions of the variables. Variable interactions, codependencies between features, result in marginal plots that do not isolate the specific contribution of a feature of interest to the target. For example, a marginal plot of sex (male/female) against body weight would likely show that, on average, men are heavier than women. While true, men are also taller than women on average, which likely accounts for most of the difference in average weight. It is unlikely that two "identical" people, differing only in sex, would be appreciably different in weight.

Rather than looking directly at the data, there are several partial dependence techniques that interrogate fitted models provided by the user: Friedman's original partial dependence [3] (which we will denote FPD), functional ANOVA [5], Individual Conditional Expectations (ICE) [4], Accumulated Local Effects (ALE) [1],

Terence Parr
University of San Francisco
E-mail: parrt@cs.usfca.edu

James D. Wilson
University of San Francisco
E-mail: jdwilson4@usfca.edu

Fig. 1: Meaningfully different results from a single partial dependence technique, FPD/ICE, applied to the same data but different models. Plots of number of bathrooms versus rent price using New York City apartment rent data [8] with $n = 10,000$ of ~50k. (a) marginal plot, (b) plot derived from random forest, (c) plot derived from gradient boosted machine, and (d) plot derived from ordinary least squares regression. Hyper parameters were tuned using 5-fold cross validation grid search over several hyper parameters. Keras model trained by experimentation: single hidden layer of 100 neurons, 500 epochs, batch size of 1000, batch normalization, and 30% dropout. STRATPD gives a plausible roughly result that rent goes up linearly with the number of bathrooms. R^2 were computed on 20% validation sets.

and most recently SHAP [11]. Model-based techniques dominate the partial dependence research literature because interpreting the output of a fitted model has several advantages. Models have a tendency to smooth over noise. Models act like analysis preprocessing steps, potentially reducing the computational burden on model-based partial dependence techniques; e.g., ALE is $O(n)$ for the n records of $\mathbf{X}$. Model-based techniques are typically model-agnostic, though for efficiency, some provide model-specific optimizations, as SHAP does. Partial dependence techniques that interrogate models also provide insight into the models themselves; i.e., how variables affect model behavior. It is also true that, in some cases, a predictive model is the primary goal so creating a suitable model is not an extra burden.

Model-based techniques do have two significant disadvantages, however. The first relates to their ability to tease apart the effect of codependent features because models are required to extrapolate into regions of nonexistent support or even into nonsensical observations; e.g., see discussions in [1] and [5]. As we demonstrate in Section 4, using synthetic and real data sets, model-based techniques can vary in their ability to isolate variable effects in practice. Second, recall that there are vast armies of business analysts and scientists at work that need to analyze data, in a manner akin to exploratory data analysis (EDA), that have no intention of creating a predictive model. Either they have no need, perhaps needing only partial dependence plots, or they do not have the expertise to choose, tune, and assess models (or write software at all).

Even in the case where a machine learning practitioner is available to create a fitted model for the analyst, hazards exist. First, if a fitted model is unable to accurately capture the relationship between features and y accurately, for whatever reason, then partial dependence does not provide any useful information to the user. To make interpretation more challenging, there is no definition of "accurate enough." Second, given an accurate fitted model, business analysts and scientists are still peering at the data through the lens of the model, which can distort partial dependence curves. Separating visual artifacts of the model from real effects present in the data requires expertise in model behavior (and optimally in the implementation of model fitting algorithms).

Consider the combined FPD/ICE plots shown in Figure 1 derived from several models (random forest, gradient boosting, linear regression, deep learning) fitted to the same New York City rent data set [8]. The subplots in Figure 1(b)-(e) present starkly different partial dependence relationships and it is unclear which, if any, is correct. The marginal plot, (a), drawn directly from the data shows a roughly linear growth in price for a rise in the number of bathrooms, but this relationship is biased because of the dependence of bathrooms on other variables, such as the number of bedrooms. (e.g., five bathroom, one bedroom apartments are unlikely.) For real data sets with codependent features, the true relationship is unknown so it is hard to evaluate the correctness of the plots. (Humans are unreliable estimators, which is why we need data analysis algorithms in the first place.) Nonetheless, having the same algorithm, operating on the same data, give meaningfully different partial dependences is undesirable and makes one question their validity.

Experts are often able to quickly recognize model artifacts, such as the stairstep phenomenon inherent to the decision tree-based methods in Figure 1(b) and (c). In this case, though, the stairstep is more accurate than the linear relationship in (d) and (e) because the number of bathrooms is discrete (except for "half baths"). The point is that interpreting model-based partial dependence plots can be misleading, even for experts.

An accurate mechanism to compute partial dependences that did not peer through fitted models would be most welcome. Such partial dependence curves would be accessible to users, like business analysts, who lack the expertise to create suitable models and would also reduce the chance of plot misinterpretation due to model artifacts. The curves could also help machine learning practitioners to choose appropriate models based upon relationships exposed in the data.

In this paper, we contribute a strategy, called STRATified Partial Dependence (STRATPD), that computes partial dependences directly from training data $(\mathbf{X}, \mathbf{y})$, rather than through the predictions of a fitted model, and does not presume mutually-independent features. As an example, Figure 1(f) shows the partial dependence plot computed by STRATPD. For numerical variables, STRATPD operates very much like a "model-free" ALE. Our technique is also based upon the notion of an idealized partial dependence: integration over the partial derivative of y with respect to the variable of interest for the smooth function that generated $(\mathbf{X}, \mathbf{y})$. As that function is unknown, we estimate the partial derivatives from the data non-parametrically. Colloquially, the approach examines changes in y across x_j while holding $x_{\setminus j}$ constant or nearly constant ($x_{\setminus j}$ denotes all variables except x_j). To hold $x_{\setminus j}$ constant, we use a single decision tree to partition feature space, a concept used by [12] and [2] for conditional permutation importance and observation similarity measures, respectively. Our second contribution, CATSTRATPD, computes partial dependence curves for categorical variables that, unlike existing techniques, does not assume adjacent category levels are similar. Both STRATPD and CATSTRATPD are quadratic in n, in the worst case (like FPD), though STRATPD behaves linearly on real data sets. Our prototype is currently limited to regression, isolates only single-variable partial dependence, and cannot identify interaction effects (as ICE can). The software is available via Python package `stratx` with source at `github`, including the code to regenerate images in this paper.

We begin by describing the proposed stratification approach in Section 2 then compare STRATPD to related (model-based) work in Section 3. In Section 4, we present partial dependence curves generated by STRATPD and CATSTRATPD on real data sets, contrast the plots with those of existing methods, and use synthetic data to highlight possible bias in some model-based methods.

2 Partial dependence without model predictions

In special circumstances, we know the precise effect of each feature x_j on y. Assume we are given training data pair $(\mathbf{X}, \mathbf{y})$ where $\mathbf{X} = [x^{(1)}, \ldots, x^{(n)}]$ is an $n \times p$ matrix whose p columns represent observed features and $\mathbf{y}$ is the $n \times 1$ vector of responses. For any smooth function $f : \mathbb{R}^p \to \mathbb{R}$ that precisely maps each $x^{(i)}$ vector to response $y^{(i)}$, $y^{(i)} = f(x^{(i)})$, the partial derivative of y with respect to x_j gives the change in y holding all other variables constant. Integrating the partial derivative then gives the *idealized partial dependence* of y on x_j, the isolated contribution of x_j to y:

Definition 1 The *idealized partial dependence* of y on feature x_j for smooth generator function $f : \mathbb{R}^p \to \mathbb{R}$ evaluated at $x_j = z$ is the cumulative sum up to z ***TODO:*** do we need to say that the partial derivative is continuous also?:

$$PD_j(z) = \int_{min(x_j)}^{z} \frac{\partial f}{\partial x_j} dx_j \qquad (1)$$

$PD_j(z)$ is the value contributed to f by x_j at $x_j = z$ and $PD_j(min(x_j)) = 0$. The advantages of this definition are that it does not depend on predictions from a fitted model and is insensitive to codependent features. Although the underlining generator function is unknown, we can estimate its partial derivatives from the raw training data. ALE also derives partial dependence by estimating and integrating across partial derivatives (e.g., see Equation 7 in [1]) but does so using local changes in fitted model predictions.

The key idea is to stratify $x_{\setminus j}$ feature space into disjoint regions of observations where all $x_{\setminus j}$ variables are approximately matched across the observations in that region. Within each $x_{\setminus j}$ region, any fluctuations in the response variable are likely due to the variable of interest, x_j. Estimates of the partial derivative within a region are estimated discretely as the changes in y values between unique and ordered x_j positions: $(y^{(i+1)} - y^{(i)})/(x_j^{(i+1)} - x_j^{(i)})$ for all i in a region. This amounts to performing piecewise linear regression through the leaf observations, one model per unique pair of x_j values, and collecting the model β_1 coefficients to estimate partial derivatives. The overall partial derivative at $x_j = z$ is the average of all slopes, found in

any region, whose range $[x_j^{(i)}, x_j^{(i+1)})$ spans z. The partial dependence curve points are, therefore, often the result of two averaging operations, which tends to smooth the curve; e.g., see Figure 3(d) below.

Stratification occurs through the use of a decision tree fit to $(\mathbf{X}_{\backslash j}, \mathbf{y})$, whose leaves aggregate observations with equal or similar $x_{\backslash j}$ features. The $x_{\backslash j}$ features can be numerical variables or label-encoded categorical variables (assigned a unique integer). STRATPD only uses the tree for the purpose of partitioning feature space and never uses predictions from any model. See Algorithm 1 for more details.

For the stratification approach to work, decision tree leaves must satisfactorily stratify $x_{\backslash j}$. If the $x_{\backslash j}$ observations in each region are not similar enough, the relationship between x_j and y is less accurate. Regions can also become so small that even the x_j values become equal, leaving a single unique x_j observation in a leaf. Without a change in x_j, no partial derivative estimate is possible and these nonsupporting observations must be ignored (e.g., the two leftmost points in Figure 2(a)). A degenerate case occurs when identical or nearly identical x_j and x_j' variables exist. Flattening x_j as part of $x_{\backslash j}$ would also flatten x_j', leading to both exhibiting flat curves, as if the decision tree were trained on $(\mathbf{X}, \mathbf{y})$ not $(\mathbf{X}_{\backslash j}, \mathbf{y})$. Our experiments show that using the collection of leaves from a random forest, which restricts the number of variables available during node splitting, prevents partitioning from relying too heavily on either x_j or x_j'. Some leaves have observations that vary in x_j or x_j' and partial derivatives can still be estimated.

STRATPD uses hyper parameter `min_samples_leaf` to control the minimum number of observations in each decision tree leaf. Generally speaking, smaller values lead to more confidence that fluctuations in y are due solely to x_j, but more observations per leaf allow STRATPD to capture more nonlinearities and make it less susceptible to noise. As the leaf size grows, however, one risks introducing contributions from $x_{\backslash j}$ into the relationship between x_j and y. At the extreme, the decision tree would consist of a single leaf node containing all observations, leading to a marginal not partial dependence curve.

STRATPD uses another hyper parameter called `min_slopes_per_x` to ignore any partial derivatives estimated with too few observations. Dropping uncertain partial derivatives greatly improves accuracy and stability. Partial dependences computed by integrating over local partial derivatives are highly sensitive to partial derivatives computed at the left edge of any x_j's range because imprecision at the left edge affects the entire curve. This presents a problem when there are few samples with x_j values at the extreme left (see, for example, the x_j histogram of Figure 8(d)). Fortunately, sensible defaults for STRATPD (10 observations and 5 slopes) work well in most cases and were used to generate all plots in this paper.

For categorical variables, CATSTRATPD uses the same stratification approach, but cannot apply regression of y to non-ordinal, categorical x_j. Instead, CATSTRATPD groups leaf observations by category and computes the average y per category in each leaf. Then, within each leaf, it chooses a random reference category and subtracts that category's $\bar{y}$ value from the leaf $\bar{\mathbf{y}}$ vector to get a vector of relative deltas between categories: $\Delta\mathbf{y} = \bar{\mathbf{y}} - \bar{\mathbf{y}}_{refcat}$. The $\Delta\mathbf{y}$ vectors from all leaves are then merged via averaging, weighted by the number of observations per category, to get the overall effect of each category on y. The delta vectors for two leaves, $\Delta\mathbf{y}$ and $\Delta\mathbf{y}'$, can only be merged if there is at least one category in common. CATSTRATPD initializes a running average vector to the first leaf's $\Delta\mathbf{y}$ and then makes passes over the remaining vectors, merging any vectors with a category in common with the running average vector. Observations associated with any remaining, unmerged leaves must be ignored.

Stratification of high-cardinality categorical variables tends to create small groups of category subsets, which complicates the averaging process across groups. (Such $\Delta\mathbf{y}$ vectors are sparse and *NaN*s represent missing values.) If both groups have the same reference category, merging is a simple matter of averaging the two delta vectors, where $mean(z, NaN) = z$. For delta vectors with different reference categories and at least one category in common, one vector is adjusted to use a randomly-selected reference category, c, in common: $\Delta\mathbf{y}' = \Delta\mathbf{y}' - \Delta\mathbf{y}_c' + \Delta\mathbf{y}_c$. That equation adjusts vector $\Delta\mathbf{y}'$ so $\Delta\mathbf{y}_c' = 0$ then adds the corresponding value from $\Delta\mathbf{y}$ so $\Delta\mathbf{y}_c' = \Delta\mathbf{y}_c$, which renders the average of $\Delta\mathbf{y}$ and $\Delta\mathbf{y}'$ meaningful. CATSTRATPD uses a single hyper parameter `min_samples_leaf` to control stratification. See Algorithm 2 for more details.

STRATPD and CATSTRATPD both have theoretical worst-case time complexity of $O(n^2)$ for n observations. For STRATPD, stratification costs $O(pn \log n)$, computing y deltas for all observations among the leaves has linear cost, and averaging slopes across unique x_j ranges is on the order of $|unique(\mathbf{X}_j)| \times n$ or n^2 when all $\mathbf{X}_j$ are unique in the worst case. STRATPD is, thus, $O(n^2)$ in the worst case. CATSTRATPD also stratifies in $O(pn \log n)$ and computes category deltas linearly in n but must make multiple passes over the $|T|$ leaves to average all possible leaf category delta vectors. In practice, three passes is the max we have seen (for high-cardinality variables), so we can assume the number of passes is some small constant. Averaging

two vectors costs $|unique(\mathbf{X}_j)|$, so each pass requires $|T| \times |unique(\mathbf{X}_j)|$. The number of leaves is roughly $n/\texttt{min\_samples\_leaf}$ and, worst-case, $|unique(\mathbf{X}_j)| = n$, meaning that merging dominates CatStratPD complexity leading to $O(n^2)$. Experiments show that our prototype is fast enough for practical use (see Section 4).

3 Related work

[3] defines the partial dependence of $\hat{f}$ on feature subsets, $S \subset F = \{1, 2, .., p\}$, as an expectation conditioned on the remaining variables:

$$FPD_S(x_S) = \mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})] = \frac{1}{n}\sum_{i=1}^{n}\hat{f}(x_S, x_{\backslash S}^{(i)}) \tag{2}$$

The Individual Conditional Expectation (ICE) plot [4] estimates the partial dependence of the prediction $\hat{f}$ on x_S, or single variable x_j, across individual observations. ICE produces a curve from the fitted model over all values of x_j while holding $x_{\backslash j}$ fixed: $\hat{f}_j^{(i)} = \widehat{f}(\{x_j^{(k)}\}_{k=1}^n, x_{\backslash j}^{(i)})$; the FPD curve for x_j is the average over all ICE curves for x_j. The motivation for ICE is to identify variable interactions that average out in the FPD curve.

SHAP [11] has roots in *Shapley regression values* [10] and calculates the average marginal effect of adding x_j to models trained on all possible subsets of features:

$$\phi_j(\hat{f}, \mathbf{x}) = \sum_{S \subseteq F \backslash \{j\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} [\hat{f}_{S \cup \{j\}}(x_{S \cup \{j\}}) - \hat{f}_S(x_S)] \tag{3}$$

To avoid training a combinatorial explosion of models with the various feature subsets, $\hat{f}_S(x_S)$, SHAP reuses a single model fitted to $(\mathbf{X}, \mathbf{y})$ by running simplified feature vectors into the model. As [13] points out, there are many possible implementations for simplified vectors. One is to replace "missing" features with their expected value or some other baseline vector (BShap in [13]). SHAP uses a more general approach ("interventional" mode) that approximates $\hat{f}_S(x_S)$ with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}'_{\backslash S})|\mathbf{X}'_S = x_S]$ where $\mathbf{X}'$ is called the *background set* and users can pass in, for example, a single vector with $\mathbf{X}_{\backslash S}$ column averages or even the entire training set, $\mathbf{X}$, which is what we will assume (and is called CES($\hat{D}$) in [13] where $\hat{D} = \mathbf{X}$). To further reduce computation costs, SHAP users typically explain a very small subsample of the data set, but with potentially a commensurate reduction in the explanatory resolution of the underlying population. SHAP has model-type-dependent optimizations for linear regression, deep learning, and decision-tree based models.

For efficiency, SHAP approximates $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})|\mathbf{X}_S = x_S]$ with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})]$, which assumes feature independence. [6] argues that "*unconditional* [as implemented] rather than *conditional* [as defined] expectations provide the right notion of dropping features." But, using the unconditional expectation makes the inner difference of Equation 3 a function of codependency-sensitive FPDs:

$$\begin{aligned}\hat{f}_{S \cup \{j\}}(x_{S \cup \{j\}}) - \hat{f}_S(x_S) &= \mathbb{E}[\hat{f}(x_{S \cup \{j\}}, \mathbf{X}_{\backslash (S \cup \{j\})})] - \mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})] \\ &= FPD_{S \cup \{j\}}(\mathbf{x}) - FPD_S(\mathbf{x})\end{aligned} \tag{4}$$

If the individual contributions are potentially biased, averaging the contributions of many such feature permutations might not lead to an accurate partial dependence. Even if the conditional expectation is used, [13] points out that SHAP is sensitive to the sparsity of $\mathbf{X}$ because condition $\mathbf{X}_S = x_S$ will find few or no training records with the exact x_S values of some input vector.

The goal of ALE ([1]) is to overcome the bias in previous model-based techniques arising from extrapolations of $\hat{f}$ far outside the support of the training data in the presence of codependent variables. ALE partitions range $[\min(\mathbf{X}_j) .. \max(\mathbf{X}_j)]$ for variable x_j into K bins and estimates the "uncentered main effect" (Equation 15) at $x_j = z$ as the cumulative sum of the partial derivatives for all bins up to the bin containing z. ALE estimates the partial derivative of $\hat{f}$ at $x_j = z$ as $\mathbb{E}[\hat{f}(b_k, \mathbf{X}_{\backslash j}) - \hat{f}(b_{k-1}, \mathbf{X}_{\backslash j}) | x_j \in (b_{k-1}, b_k]]$ for bin b_k that contains z. They also extend ALE to two variables by partitioning feature space into K^2 rectangular bins and computing the second-order finite difference of $\hat{f}$ with respect to the two variables for each bin.

Another related technique that integrates over partial derivatives to measure x_j effects is called Integrated Gradients (IG) from [14]. Given a single input vector, $\mathbf{x}$, to a deep learning network, IG integrates over the

loss function gradient at points along the path from a baseline vector, $\mathbf{x}'$, to $\mathbf{x}$. IG can be seen as computing the partial dependence of the loss function, rather than model output, $\hat{f}$, on the x_j values in $\mathbf{x}$.


Fig. 2: Comparison of STRATPD, CATSTRATPD and ALE for $p = 2$ with continuous x_2 and continuous x_1 in (a), (b) and categorical x_1 in (c), (d). Vertical dashed lines are ALE bin edges and horizontal dashed lines are regions partitioned by a decision tree fit to $(\mathbf{X}_{\backslash j}, \mathbf{y})$. Whereas ALE shifts x_1 observations to bin edges holding x_2 exactly constant (and asks for predictions at those points), STRATPD assumes x_2 values are the same and uses known y values. The small wedges on x_1 axis indicate partial dependence curve points. The shades of green and blue indicate y values. The leftmost two observations in (a) and the lowermost observation in (c) are ignored as finite differences are not defined.

STRATPD is like a "model-free" version of ALE in that STRATPD also defines partial dependence as the cumulative sum of partial derivatives, but we estimate derivatives using $\mathbf{y}$ values rather than $\hat{f}$ predictions. ALE partitions x_j into bins then fixes $x_{\backslash j}$ as it shifts x_j to bin edges to compute finite differences, as depicted in Figure 2(b) for $p = 2$. The wedges on the x_1 axis indicate the x_1 points of the computed partial dependence curve. STRATPD partitions $x_{\backslash j}$ into regions of (hopefully) similar observations and computes finite differences between the average y values at unique x_j values in each region, as depicted in Figure 2(a). The leftmost two observations are ignored as there is no change in x_1 in that leaf. The shaded area illustrates that the partial derivative at any $x_j = z$ is the average of all derivatives spanning z across $x_{\backslash j}$ regions. STRATPD assumes all points within a region are identical in $x_{\backslash j}$, effectively projecting points in $x_{\backslash j}$ space onto a hyperplane if they are not. ***TODO:*** is projection right term? ALE shifts x_j values in a small neighborhood and STRATPD depends on a suitable `min_samples_leaf` hyper parameter to prevent $x_{\backslash j}$ points in a regions from becoming too dissimilar. STRATPD automatically generates more curve points in areas of high x_j density, but ALE is more efficient.

Using decision trees for the purpose of partitioning feature space as STRATPD does was previously used by [12] to improve permutation importance for random forests by permuting x_j only within the observations of a leaf. Earlier, [2] defined a similarity measure between two observations according to how often they appear in the same leaf in a random forest. Rather than partitioning $x_{\backslash j}$ space like STRATPD, those techniques partitioned all of x space.

The model-based techniques under discussion treat boolean and label-encoded categorical variables (encoded as unique integers) as numerical variables, even though there is no defined order, as depicted in Figure 2(d). ALE does, however, take advantage of the lack of order to choose an x_j order that reduces "extrapolation outside the data envelope" by measuring the similarity of $\mathbf{X}_{\backslash j}$ sample values across x_j categories. Adjacent category integers, though, could still represent the most semantically different categories, so any shift of an observation's category to extrapolate is risky. Even the smallest possible extrapolation can conjure up nonsensical observations, such as pregnant males, as we demonstrate in Figure 7 below where FPD, SHAP, and ALE underestimate pregnancy's effect on body weight. (For boolean x_j, ALE behaves like FPD.)

In contrast, CATSTRATPD uses a different algorithm for categoricals and computes differences between the average y for all categories within the leaf to a random reference category; see Figure 2(c). These leaf delta vectors are then merged across leaves to arrive at an overall delta vector relating the relative effect of each category on y. One could argue that STRATPD also extrapolates because $x_{\backslash j}$ could include categorical variables and not all $x_{\backslash j}$ records would be identical. But, our approach only assumes $x_{\backslash j}$ values are similar and uses known training y values for finite differences, rather than asking a model to make prediction for nonsensical records, which could be wildly inaccurate. Also, the decision tree would, by definition, likely partition $x_{\backslash j}$ space into regions that can be treated similarly, thus, grouping semantically similar categorical levels together.

4 Experimental results

In this section, we demonstrate experimentally that STRATPD and CATSTRATPD compute accurate partial dependence curves for synthetic data and plausible results for a real data set. Experiments also provide evidence that existing model-based techniques can provide meaningfully-biased curves. We begin by comparing the partial dependence curves from popular techniques on synthetic data with complex interactions.[1]


Fig. 3: Partial dependence plots of $n = 2000$ data generated from noiseless $y = x_1^2 + x_1 x_2 + 5x_1 sin(3x_2) + 10$ where $x_1, x_2, x_3 \sim U(0, 10)$ and x_3 does not affect y. The model is a random forest with 10 trees getting training $R^2 = 0.997$. SHAP used all $\mathbf{X}$ as background data and ALE used $K = 300$. The curves were generated from same 2000 data points that the model was trained on.

Figure 3 illustrates that FPD, SHAP, ALE, and STRATPD all suitably isolate the effect of independent individual variables on the response for noiseless data generated via: $y = x_1^2 + x_1 x_2 + 5x_1 sin(3x_2) + 10$ for $x_1, x_2, x_3 \sim U(0, 10)$ where x_3 does not affect y. The shapes of the curves for all techniques look similar except that STRATPD starts all curves at $y = 0$ (as could the others). SHAP's curves have the advantage that they indicate the presence of variable interactions. To our eye, STRATPD's curves are smoothest despite not having access to model predictions.

Models have a tendency to smooth out noise and a legitimate concern is that, without the benefit of a model, STRATPD could be adversely affected. Figure 4 demonstrates STRATPD curves for noisy quadratics generated from $y = x_1^2 + x_2 + 10 + \epsilon$ where $\epsilon \sim N(0, \sigma)$ ϵ and, at $\sigma = 2$, 95% of the noise falls within [0,4] (since $2\sigma = 4$), meaning that the signal-to-noise ratio is at best 1-to-1 for x_1^2 in $[-2, 2]$. For zero-centered Gaussian noise and this data set, STRATPD accuracy is stable. The superfluous noise variable x_3 in Figure 3 also did not confuse STRATPD.


Fig. 4: The effective noise on STRATPD for $y = x_1^2 + x_1 + 10 + \epsilon$ where $x_1, x_2 \sim U(-2, 2)$, $\epsilon \sim N(0, \sigma)$ with $\sigma \in \{0, 0.5, 1, 2\}$. The 95% interval for amplitude of the noise in (d) for $\sigma = 2$ is the same as the signal.

Turning to categorical variables, Figure 5 presents partial dependence curves for FPD, ALE, and CAT-STRATPD derived from a noisy synthetic weather data set, where temperature varies in sinusoidal fashion over the year and with different baseline temperatures per state. (The vertical "smear" in the FPD plot shows the complete sine waves but from the side, edge on.) Variable `state` is independent and all plots identify the baseline temperature per state correctly.

The primary goal of the stratification approach proposed in this paper is to obtain accurate partial dependence curves in the presence of codependent variables. To test STRATPD/CATSTRATPD and discover

Fig. 5: $y = base[x_{state}] + 10sin(\frac{2\pi}{365}x_{dayofyear} + \pi) + \epsilon$ where $\epsilon \sim N(\mu = 0, \sigma = 4)$. The *base* temperature per state is $\{AZ = 90, CA = 70, CO = 40, NV = 80, WA = 60\}$. Sinusoids in (a) are the average of three years' temperature data.

potential bias in existing techniques, we synthesized a body weight data set generated by the following equation with nontrivial codependencies between variables:

$$
\begin{aligned}
y = \quad & 120 + 10(x_{height} - min(x_{height})) + 40x_{pregnant} - 1.5x_{education} \\
\text{where } & x_{sex} \sim Bernoulli(\{M, F\}, p = 0.5) \\
& x_{pregnant} = \begin{cases} Bernoulli(\{0, 1\}, p = 0.5) & \text{if } x_{sex} = F \\ 0 & \text{if } x_{sex} = M \end{cases} \\
& x_{height} = \begin{cases} 5 * 12 + 5 + \epsilon & \text{if } x_{sex} = F, \ \epsilon \sim U(-4.5, 5) \\ 5 * 12 + 8 + \epsilon & \text{if } x_{sex} = M, \ \epsilon \sim U(-7, 8) \end{cases} \\
& x_{education} = \begin{cases} 12 + \epsilon & \text{if } x_{sex} = F, \ \epsilon \sim U(0, 8) \\ 10 + \epsilon & \text{if } x_{sex} = M, \ \epsilon \sim U(0, 8) \end{cases}
\end{aligned}
\tag{5}
$$

The partial derivative of y with respect to x_{height} is 10 (holding all other variables constant), so the optimal partial dependence curve is a line with slope 10. Figure 6 illustrates the curves for the techniques under consideration, with ALE and StratPD giving the sharpest representation of the linear relationship. (StratPD's curve is drawn on top of the SHAP plots using the righthand scale.) The FPD and both SHAP plots also suggest a linear relationship, albeit with a little less precision. The ICE curves in Figure 6(a) and "fuzzy" SHAP curves have the advantage that they alert users to variable dependencies or interaction terms. On the other hand, the kink in the partial dependence curve and other visual phenomena could confuse less experienced machine learning practitioners and certainly analysts and researchers in other fields (our primary target communities).


Fig. 6: Partial dependence plots of response body weight on feature x_{height} using n=2000 synthetic observations from Equation 5. FPD and SHAP have "kinks" at the maximum female height. SHAP defines feature importance as the average SHAP value magnitude, which overemphasizes importance for heights below 70 inches here. The male/female ratio is 50/50, half of the women are pregnant, and pregnancy contributes 40 pounds. SHAP interrogated an RF tuned via 5-fold cross validation grid search (OOB R^2 0.999) and explained all 2000 samples; the interventional case used 100 observations as background data.

Even for experts, explaining this behavior requires some thought, and one must distinguish between model artifacts and interesting phenomena. The discontinuity at the maximum female height location arises partially from the model having trouble extrapolating for extremely tall pregnant women. Consider one of the upper ICE lines in Figure 6(a) for a pregnant woman. As the ICE line slides x_{height} above the maximum height for a woman, the model leaves the support of the training data and predicts a *lower* weight as height increases. ALE's curve is straight because it focuses on local effects, demonstrating that the lack of sharp slope-10 lines for FPD and SHAP cannot be attributed simply to a poor choice of model.

SHAP also defines x_j feature importance as the average magnitude of the x_j SHAP values, which introduces a paradox. The spread of the SHAP values alerts users to variable interactions, but allows contributions from other variables to leak in, thus, potentially leading to less precise estimates of x_{height}'s importance.

Fig. 7: Partial dependence bar charts for boolean $x_{pregnant}$ with the same model from Figure 6. Only CatStratPD gets the correct 40 pound contribution per Equation 5.

It is conceivable that a more sophisticated model (in terms of extrapolation) could sharpen the FPD and SHAP curves for x_{height}. There is a difference, however, between extrapolating to a meaningful but unsupported vector and making predictions for nonsensical vectors arising from variable codependencies. Techniques that rely on such predictions make the implicit assumption of variable independence, introducing the potential for bias. Consider Figure 7 that presents the partial dependence results for categorical variable $x_{pregnant}$ (same data set). The weight gain from pregnancy is 40 pounds per Equation 5, but only CatStratPD identifies that exact relationship; FPD, SHAP, and ALE show a gain of 30 pounds.

CatStratPD stratifies persons with the same or similar sex, education, and height into groups and then examines the relationship between $x_{pregnant}$ and y. If a group contains both pregnant and nonpregnant females, the difference in weight will be 40 pounds in this noiseless data set (if we assume identical $x_{\setminus j}$). FPD and ALE rely on computations that require fitted models to conjure up predictions for nonsensical records representing pregnant males (e.g., $\hat{f}(x_j = pregnant, \mathbf{X}_{\setminus j})$). Not even a human knows how to estimate the weight of a pregnant male. SHAP, per its definition, does not require such predictions, but in practice for efficiency reasons, SHAP approximates $\mathbb{E}[\hat{f}(x_j = pregnant, \mathbf{X}_{\setminus j})|\mathbf{X}_j = x_j]$ with $\mathbb{E}[\hat{f}(x_j = pregnant, \mathbf{X}_{\setminus j})]$, which does not restrict pregnancy to females. As discussed above, there are advantages to all of these model-based techniques, but this example demonstrates there is potential for partial dependence bias.


Fig. 8: Partial dependence curves of bulldozer `YearMade` versus `SalePrice` for FPD/ SHAP, ALE, and StratPD. n=10,000 observations drawn from ~362k. SHAP interrogated a random forest (OOB $R^2 = 0.808$) to explain 1000 training observations with 100 observations as background data. Hyper parameters were tuned using 5-fold cross validation grid search over several hyper parameters.

The stratification approach also gives plausible results for real data sets, such as the bulldozer auction data from [9]. Figure 8 shows the partial dependence curves for the same set of techniques as before on feature `YearMade`, chosen because it is very predictive of sale price. The shape and magnitude of the FPD, SHAP, ALE, and StratPD curves are similar, indicating that older bulldozers are worth less at auction, which is plausible. The StratPD curve shows 10 bootstrapped trials where the heavy black dots represent the partial dependence curve and the other colored curves describe the variability.

And, finally, an important consideration for any tool is performance, so we plotted execution time versus data size (up to 30,000 observations) for three real Kaggle data sets: rent, bulldozer, and flight arrival delays [7]. Figure 9 shows growth curves for 40 numerical variables and 11 categorical variables grouped by type of variable. For these data sets, StratPD takes 1.2s or less to process 30,000 records for any numerical x_j, despite the potential for quadratic cost. CatStratPD typically processes categorical variables in less than 2s but takes 13s for the high-cardinality categorical `ModelID` of bulldozer (which looks mildly quadratic). These elapsed times for our prototype show it to be practical and competitive with FPD/ICE, SHAP, and ALE. If the cost to train a model using (cross validated) grid search for hyper parameter tuning is included, StratPD and CatStratPD outperforms these existing techniques (as training and tuning is often measured in minutes).

Fig. 9: Time to compute partial dependence curve for up to 30,000 observations for 40 numerical and 11 categorical variables timing blah. $p = 20$, bulldozer $p = 14$, and flight $p = 17$. The Numba just-in-time compiler is used to improve performance, but timing does not include compilation; users do experience this "warm-up" time.

5 Conclusion and future work

In this paper, we contribute a method for computing partial dependence curves, for both numerical and categorical variables, that does not use predictions from a provided fitted model. Working directly from the data makes partial dependences accessible to business analysts and scientists not qualified to choose, tune, and assess machine learning models. For experts, it can provide hints about the relationships in the data to help guide their choice of model. Our experiments show that StratPD and CatStratPD are fast enough for practical use and correctly identify partial dependences for synthetic data and give plausible curves on real data sets. StratPD relies on two important hyper parameters (with broadly applicable defaults) but model-based techniques should include the hyper parameters of the required fitted model for a fair comparison. Our goal here is not to argue that model-based techniques are not useful. Rather, we are pointing out potential issues and hoping to open a new line of nonparametric inquiry that experiments have shown to be applicable in situations and accurate in cases where model-based techniques are not. The obvious next goal for this approach is a version for classifiers and to extend the technique to two variables, paralleling ALE's second order derivative approach.

References

1. Apley, D.W., Zhu, J.: Visualizing the effects of predictor variables in black box supervised learning models. arXiv preprint arXiv:1612.08468 (2016)
2. Breiman, L., Cutler, A.: Random forests website. https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm (2003). Accessed: 2019-05-24
3. Friedman, J.H.: Greedy function approximation: A gradient boosting machine. Annals of Statistics **29**, 1189–1232 (2000)
4. Goldstein, A., Kapelner, A., Bleich, J., Pitkin, E.: Peeking inside the black box: Visualizing statistical learning with plots of individual conditional expectation. Journal of Computational and Graphical Statistics **24**(1), 44–65 (2015). DOI 10.1080/10618600.2014.907095. URL https://doi.org/10.1080/10618600.2014.907095
5. Hooker, G.: Generalized functional anova diagnostics for high-dimensional functions of dependent variables. Journal of Computational and Graphical Statistics **16**(3), 709–732 (2007). DOI 10.1198/106186007X237892. URL https://doi.org/10.1198/106186007X237892
6. Janzing, D., Minorics, L., Blöbaum, P.: Feature relevance quantification in explainable ai: A causal problem (2019)
7. Kaggle: 2015 flight delays and cancellations. https://www.kaggle.com/usdot/flight-delays (2015). Accessed: 2019-12-01
8. Kaggle: Two sigma connect: Rental listing inquiries. https://www.kaggle.com/c/two-sigma-connect-rental-listing-inquiries (2017). Accessed: 2019-06-15
9. Kaggle: Blue book for bulldozer. https://www.kaggle.com/sureshsubramaniam/blue-book-for-bulldozer-kaggle-competition (2018). Accessed: 2019-06-15
10. Lipovetsky, S., Conklin, M.: Analysis of regression in game theory approach. Applied Stochastic Models in Business and Industry **17**, 319 – 330 (2001). DOI 10.1002/asmb.446
11. Lundberg, S.M., Lee, S.I.: A unified approach to interpreting model predictions. In: I. Guyon, U.V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, R. Garnett (eds.) Advances in Neural Information Processing Systems 30, pp. 4765–4774. Curran Associates, Inc. (2017). URL http://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions.pdf
12. Strobl, C., Boulesteix, A.L., Kneib, T., Augustin, T., Zeileis, A.: Conditional variable importance for random forests. BMC Bioinformatics **9**(1), 307 (2008). DOI 10.1186/1471-2105-9-307. URL https://doi.org/10.1186/1471-2105-9-307
13. Sundararajan, M., Najmi, A.: The many shapley values for model explanation. ArXiv **abs**/**1908.08474** (2019)
14. Sundararajan, M., Taly, A., Yan, Q.: Axiomatic attribution for deep networks. In: Proceedings of the 34th International Conference on Machine Learning - Volume 70, ICML'17, p. 3319–3328. JMLR.org (2017)

6 Appendix

Algorithm 1: $StratPD(\mathbf{X}, \mathbf{y}, j, min\_samples\_leaf, min\_slopes\_per\_x)$

$T :=$ Decision tree regressor fit to $(\mathbf{X}_{\setminus j}, \mathbf{y})$ with hyper parameter: $min\_samples\_leaf$

for *each leaf* $L \in T$ **do**

    $(\mathbf{x}_L, \mathbf{y}_L) := \{(x_j^{(i)}, y^{(i)})\}_{i \in L}$ $\triangleright$ *Get leaf observations*

    $\mathbf{ux}, \bar{\mathbf{y}} :=$ Group and sort $(\mathbf{x}_L, \mathbf{y}_L)$ by x_j value, computing $\bar{y}$ per unique x_j value

    $\mathbf{dx} := \mathbf{ux}^{(i+1)} - \mathbf{ux}^{(i)}$ for $i = 1..|\mathbf{ux}| - 1$ $\triangleright$ *Discrete difference between adjacent*

    $\mathbf{dy} := \bar{\mathbf{y}}^{(i+1)} - \bar{\mathbf{y}}^{(i)}$ for $i = 1..|\mathbf{ux}| - 1$

    Add tuples $(\mathbf{ux}^{(i)}, \mathbf{ux}^{(i+1)}, \mathbf{dy}^{(i)}/\mathbf{dx}^{(i)})$ to list $\mathbf{D}$ for $i = 1..|\mathbf{ux}| - 1$

$\mathbf{ux} := unique(\mathbf{X}_j)$

for *each* $x \in \mathbf{ux}$ **do** $\triangleright$ *Count slopes and compute average slope per unique x_j value*

    $\mathbf{slopes} := [slope$ for $(a, b, slope) \in \mathbf{D}$ if $x \geq a$ and $x < b]$

    $\mathbf{c}_x := |\mathbf{slopes}|$

    $\mathbf{dydx}_x := \overline{\mathbf{slopes}}$

$\mathbf{dydx} := \mathbf{dydx}[\mathbf{c} \geq min\_slopes\_per\_x]$ $\triangleright$ *Drop missing slopes, those computed with too few*

$\mathbf{ux} := \mathbf{ux}[\mathbf{c} \geq min\_slopes\_per\_x]$

$\mathbf{pdx} := \mathbf{ux}^{(i+1)} - \mathbf{ux}^{(i)}$ for $i = 1..|\mathbf{ux}| - 1$

$\mathbf{pdy} := [0] +$ cumulative$\_$sum$(\mathbf{dydx} * \mathbf{pdx})$ $\triangleright$ *integrate, inserting 0 for leftmost x_j*

return $\mathbf{pdx}, \mathbf{pdy}$

Algorithm 2: $CatStratPD(\mathbf{X}, \mathbf{y}, j, min\_samples\_leaf)$

$T :=$ Decision tree regressor fit to $(\mathbf{X}_{\setminus j}, \mathbf{y})$ with hyper parameter: $min\_samples\_leaf$

Let ΔY be a matrix whose columns are vectors of leaf category deltas

Let C be a matrix whose columns are vectors for leaf category counts

for *each leaf* $L \in T$ **do** $\triangleright$ *Get average y delta relative to random ref category for obs. in leaves*

    $(\mathbf{x}_L, \mathbf{y}_L) := \{(x_j^{(i)}, y^{(i)})\}_{i \in L}$ $\triangleright$ *Get leaf observations*

    $\mathbf{ucats}, C_L := unique(\mathbf{x}_L)$ $\triangleright$ *Get unique categories, counts from leaf observations*

    $\bar{\mathbf{y}} :=$ Group leaf records $(\mathbf{x}_L, \mathbf{y}_L)$ by category, computing $\bar{y}$ per unique category

    $refcat :=$ random category from $\mathbf{x}_L$

    $\Delta Y_L = \bar{\mathbf{y}} - \bar{\mathbf{y}}_{refcat}$

$\Delta\mathbf{y}, \mathbf{c} := \Delta Y_1, C_1$ $\triangleright$ *$\Delta\mathbf{y}$ is running average vector mapping category to average y delta*

$completed := \{1\}; work := \{2..|leaves|\};$ $\triangleright$ *set of leaf indexes*

while $|work| > 0$ *and* $|completed| > 0$ **do** $\triangleright$ *2 passes is typical to merge all ΔY_L into $\Delta\mathbf{y}$*

    $completed := \emptyset$

    for *each leaf* L *in work* **do**

        if ΔY_L *has category in common with* $\Delta\mathbf{y}$ **then**

            $completed := completed \cup \{L\}$

            $c =$ random category in intersection

            $\Delta\mathbf{y}_L := \Delta Y_L - \Delta Y_{L,c} + \Delta\mathbf{y}_c$ $\triangleright$ *Adjust so $\Delta Y_{L,c} = 0$, add corresponding $\Delta\mathbf{y}_c$ value*

            $\Delta\mathbf{y} := (\mathbf{c} \times \Delta\mathbf{y} + C_L \times \Delta\mathbf{y}_L)/(\mathbf{c} + C_L)$ where $z + NaN = z$ $\triangleright$ *weighted average*

            $\mathbf{c} := \mathbf{c} + C_L$ $\triangleright$ *update running weight*

    $work := work \backslash completed$

return $\Delta\mathbf{y}$