Technical Report: Partial Dependence without Model Predictions through Stratification

Terence Parr
University of San Francisco
parrt@cs.usfca.edu

James D. Wilson
University of San Francisco
jdwilson4@usfca.edu

Abstract

1 Introduction

Partial dependence, the isolated effect of a specific variable or variables on the response variable, y, is important to researchers and practitioners in many disparate fields such as medicine, business, and the social sciences. For example, in medicine, researchers are interested in the relationship between an individual's demographics or clinical features and their susceptibility to illness. Business analysts at a car manufacturer might need to know how changes in their supply chain are affecting defect rates. Climate scientists are interested in how different atmospheric carbon levels affect temperature.

For an explanatory matrix, $\mathbf{X}$, with a single variable, x_1, a plot of the y against x_1 visualizes the marginal effect of feature x_1 on y exactly. Given two or more features, one can similarly plot the marginal effects of each feature separately, however, the analysis is complicated by the interactions of the variables. Variable interactions, codependencies between features, result in marginal plots that do not isolate the specific contribution of a feature of interest to the target. For example, a marginal plot of sex (male/female) against body weight would likely show that, on average, men are heavier than women. While true, men are also taller than women on average, which likely accounts for most of the difference in average weight. It is unlikely that two "identical" people, differing only in sex, would be appreciably different in weight.

Rather than looking directly at the data, there are several partial dependence techniques that interrogate fitted models provided by the user: Friedman's original partial dependence (which we will denote FPD) Friedman [2000], functional ANOVA Hooker [2007], Individual Conditional Expectations (ICE) Goldstein et al. [2015], Accumulated Local Effects (ALE) Apley [2016], and most recently SHAP Lundberg and Lee [2017]. Model-based techniques dominate the partial dependence research literature because interpreting the output of a fitted model has several advantages. Models have a tendency to smooth over noise. Models act like analysis preprocessing steps, potentially reducing the computational burden on model-based partial dependence techniques; e.g., ALE is $O(n)$ for the n records of $\mathbf{X}$. Model-based techniques are typically model-agnostic, though for efficiency, some provide model-specific optimizations, as SHAP does. Partial dependence techniques that interrogate models also provide insight into the models themselves; i.e., how variables affect model behavior. It is also true that, in some cases, a predictive model is the primary goal so creating a suitable model is not an extra burden.

Model-based techniques do have two significant disadvantages, however. The first relates to their ability to tease apart the effect of codependent features because models are required

Figure 1: Plots of bathrooms versus rent price using New York City apartment rent data. (a) marginal plot, (b) PD/ICE plot derived from random forest, (c) PD/ICE plot derived from gradient boosted machine, and (d) PD/ICE plot derived from ordinary least squares regression; sample size is 10,000 observations of ~50k. The PD/ICE plots are different for the same data set, depending on the chosen user model. X['bathrooms'].unique() shows (array([0. , 1. , 1.5, 2. , 2.5, 3. , 3.5, 4.]), array([54, 8151, 140, 1539, 39, 67, 3, 7])). STRATPD has missing last value, not enough data. what are R^2 values. how tuned all but last share y with left

to extrapolate into regions of sparse or nonexistent support; e.g., see discussions in Apley [2016] and Hooker [2007]. As we demonstrate in Section 4, using synthetic and real data sets, model-based techniques can vary in their ability to isolated variable effects in practice. Second, recall that there are vast armies of business analysts and scientists at work that need to analyze data, in a manner akin to exploratory data analysis (EDA), that have no intention of creating a predictive model. Either they have no need, perhaps needing only partial dependence plots, or they do not have the expertise to choose, tune, and assess models (or write software at all).

Even in the case where a machine learning practitioner is available to create a fitted model for the analyst, hazards exist. First, if a fitted model is unable to accurately capture the relationship between features and y accurately, for whatever reason, then partial dependence does not provide any useful information to the user. To make interpretation more challenging, there is no definition of "accurate enough." Second, given an accurate fitted model, business analysts and scientists are still peering at the data through the lens of the model, which can distort partial dependence curves. Separating visual artifacts of the model from real effects present in the data requires expertise in model behavior (and optimally in the implementation of model fitting algorithms).

Consider the combined FPD/ICE plots shown in Figure 1 derived from several models (random forest, gradient boosting, linear regression, deep learning) fitted to the same New York City rent data set Kaggle [2017]. The subplots in Figure 1(b)-(e) present starkly different partial dependence relationships and it is unclear which, if any, is correct. The marginal plot, (a), drawn directly from the data shows a roughly linear growth in price for a rise in the number of bathrooms, but this relationship is biased because of the dependence of bathrooms on other variables, such as the number of bedrooms. (e.g., five bathroom, one bedroom apartments are unlikely.) For real data sets with codependent features, the true relationship is unknown so it is hard to evaluate the correctness of the plots. (Humans are unreliable estimators, which is why we need data analysis algorithms in the first place.) Nonetheless, having the same algorithm, operating on the same data, give meaningfully different partial dependences is undesirable and makes one question their validity.

Experts are often able to quickly recognize model artifacts, such as the stairstep phenomenon inherent to the decision tree-based methods in Figure 1(b) and (c). In this case, though, the stairstep is more accurate than the linear relationship in (d) and (e) because the number of bathrooms is discrete (except for "half baths"). The point is that interpreting model-based partial dependence plots can be misleading, even for experts.

An accurate mechanism to compute partial dependences that did not peer through fitted models would be most welcome. Such partial dependence curves would be accessible to users, like business analysts, who lack the expertise to create suitable models and would also reduce the chance of plot misinterpretation due to model artifacts. The curves could also help machine learning practitioners to choose appropriate models based upon relationships exposed in the data.

In this paper, we propose a strategy, called STRATified Partial Dependence (STRATPD), that (*i*) computes partial dependences directly from training data $(\mathbf{X}, \mathbf{y})$, rather than through the predictions of a fitted model, and (*ii*) does not presume mutually-independent features. As an example, Figure 1(f) shows the partial dependence plot computed by STRATPD. The technique depends on the notion of an idealized partial dependence: integration over the partial derivative of y with respect to the variable of interest for the smooth function that generated $(\mathbf{X}, \mathbf{y})$. As that function is unknown, we estimate the partial derivatives from the data non-parametrically. Colloquially, the approach examines changes in y across x_j while holding $x_{\setminus j}$ constant or nearly constant ($x_{\setminus j}$ denotes all observations with all variables except x_j). A similar stratification approach works for categorical variables (CATSTRATPD). Both STRATPD and CATSTRATPD are quadratic in n, in the worst case (like FPD), but STRATPD behaves linearly on real data sets. Our prototype is currently limited to regression, isolates only single-variable partial dependence, and cannot identify interaction effects (as ICE can). The software is available via Python package `stratx` with source code at `github`, including the code to regenerate images in this paper.

We begin by describing the proposed stratification approach in Section 2 then compare STRATPD to related (model-based) work in Section 3. In Section 4, we present partial dependence curves generated by STRATPD and CATSTRATPD on real data sets, contrast the plots with those of existing methods, and use synthetic data to highlight possible bias in some model-based methods.

2 Partial dependence without model predictions

In special circumstances, we know the precise effect of each feature x_j on y. Assume we are given training data pair $(\mathbf{X}, \mathbf{y})$ where $\mathbf{X} = [x^{(1)}, \ldots, x^{(n)}]$ is an $n \times p$ matrix whose p columns represent observed features and $\mathbf{y}$ is the $n \times 1$ vector of responses. For any smooth function $f : \mathbb{R}^p \to \mathbb{R}$ that precisely maps each $x^{(i)}$ vector to response $y^{(i)}$, $y^{(i)} = f(x^{(i)})$, the partial derivative of y with respect to x_j gives the change in y holding all other variables constant. Integrating the partial derivative then gives the *idealized partial dependence* of y on x_j, the isolated contribution of x_j to y:

Definition 1 The *idealized partial dependence* of y on feature x_j for smooth generator function $f : \mathbb{R}^p \to \mathbb{R}$ evaluated at $x_j = z$ is the cumulative sum up to z ***TODO:*** do we need to say that the partial derivative is continuous also?:

$$PD_j(z) = \int_{min(x_j)}^{z} \frac{\partial f}{\partial x_j} dx_j \tag{1}$$

$PD_j(z)$ is the value contributed to f by x_j at $x_j = z$ and $PD_j(min(x_j)) = 0$. The advantages of this definition are that it does not depend on predictions from a fitted model and is insensitive to codependent features. Although the underlining generator function is unknown, we can estimate its partial derivatives from the raw training data. ALE also derives partial dependence by estimating and integrating across partial derivatives (e.g., see Equation 7 in Apley [2016]) but does so using local changes in fitted model predictions.

The key idea is to stratify $x_{\setminus j}$ feature space into disjoint regions of observations where all $x_{\setminus j}$ variables are approximately matched across the observations in that region. Within each $x_{\setminus j}$ region, any fluctuations in the response variable are likely due to the variable of interest, x_j. Estimates of the partial derivative within a region are estimated discretely as the changes in y values between unique and ordered x_j positions: $(y^{(i+1)} - y^{(i)})/(x_j^{(i+1)} - x_j^{(i)})$ for all i in a region. This amounts to performing piecewise linear regression through the leaf observations, one model per unique pair of x_j values, and collecting the model β_1 coefficients to estimate partial derivatives. The overall partial derivative at $x_j = z$ is the average of all slopes, found in any region, whose range $[x_j^{(i)}, x_j^{(i+1)})$ spans z. Stratification occurs through the use of a decision tree fit to $(\mathbf{X}_{\setminus j}, \mathbf{y})$, whose leaves aggregate observations with equal or similar $x_{\setminus j}$ features. The $x_{\setminus j}$ features can be numerical variables or label-encoded categorical variables (assigned a unique integer). STRATPD only uses the tree for the purpose of partitioning feature space and never uses predictions from any model. See Algorithm 1 for more details.

For this approach to work, decision tree leaves must satisfactorily stratify $x_{\backslash j}$. If the $x_{\backslash j}$ observations in each region are not similar enough, the relationship between x_j and y is less accurate. Regions can also become so small that even the x_j values become equal, leaving a single unique x_j observation in a leaf. Without a change in x_j, no partial derivative estimate is possible and these nonsupporting observations must be ignored. A degenerate case occurs when identical or nearly identical x_j and x'_j variables exist. Flattening x_j as part of $x_{\backslash j}$ would also flatten x'_j, leading to both exhibiting flat curves, as if the decision tree were trained on $(\mathbf{X}, \mathbf{y})$ not $(\mathbf{X}_{\backslash j}, \mathbf{y})$. Our experiments show that using the collection of leaves from a random forest, which restricts the number of variables available during node splitting, prevents partitioning from relying too heavily on either x_j or x'_j. Some leaves have observations that vary in x_j or x'_j and partial derivatives can still be estimated. ***TODO:*** maybe talk about how PD/ICE on RF underestimates the curve.

StratPD uses hyper parameter `min_samples_leaf` to control the minimum number of observations in each decision tree leaf. Generally speaking, smaller values lead to more confidence that fluctuations in y are due solely to x_j, but more observations per leaf allow StratPD to capture more nonlinearities and make it less susceptible to noise. As the leaf size grows, however, one risks introducing contributions from $x_{\backslash j}$ into the relationship between x_j and y. At the extreme, the decision tree would consist of a single leaf node containing all observations, leading to a marginal not partial dependence curve.

StratPD uses another hyper parameter called `min_slopes_per_x` to ignore any partial derivatives estimated with too few observations. Dropping uncertain partial derivatives greatly improves accuracy and stability. Partial dependences computed by integrating over local partial derivatives are highly sensitive to partial derivatives computed at the left edge of any x_j's range because imprecision at the left edge affects the entire curve. This presents a problem when there are few samples with x_j values at the extreme left (see, for example, the x_j histogram of Figure 8(d)). Fortunately, sensible defaults for StratPD (10 observations and 5 slopes) work well in most cases and were used to generate all plots in this paper.

For categorical variables, CatStratPD uses the same stratification approach, but cannot apply regression of y to non-ordinal, categorical x_j. Instead, CatStratPD groups leaf observations by category and computes the average $\bar{y}$ per category. Then, within each leaf, it chooses a random reference category and subtracts that category's $\bar{y}$ value from the leaf $\bar{\mathbf{y}}$ vector to get a vector of relative deltas between categories: $\Delta\mathbf{y} = \bar{\mathbf{y}} - \bar{\mathbf{y}}_{refcat}$. The $\Delta\mathbf{y}$ vectors from all leaves are then merged via averaging, weighted by the number of observations per category, to get the overall effect of each category on y. The delta vectors for two leaves, $\Delta\mathbf{y}$ and $\Delta\mathbf{y}'$, can only be merged if there is at least one category in common. CatStratPD initializes a running average vector to the first leaf's $\Delta\mathbf{y}$ and then makes passes over the remaining vectors, merging any vectors with a category in common with the running average vector. Observations associated with any remaining, unmerged leaves must be ignored.

Stratification of high-cardinality categorical variables tends to create small groups of category subsets, which complicates the averaging process across groups. (Such $\Delta\mathbf{y}$ vectors are sparse and *NaN*s represent missing values.) If both groups have the same reference category, merging is a simple matter of averaging the two delta vectors, where $mean(z, NaN) = z$. For delta vectors with different reference categories and at least one category in common, one vector is adjusted to use a randomly-selected reference category, c, in common: $\Delta\mathbf{y}' = \Delta\mathbf{y}' - \Delta\mathbf{y}'_c + \Delta\mathbf{y}_c$. That equation adjusts vector $\Delta\mathbf{y}'$ so $\Delta\mathbf{y}'_c = 0$ then adds the corresponding value from $\Delta\mathbf{y}$ so $\Delta\mathbf{y}'_c = \Delta\mathbf{y}_c$, which renders the average of $\Delta\mathbf{y}$ and $\Delta\mathbf{y}'$ meaningful. CatStratPD uses a single hyper parameter `min_samples_leaf` to control stratification. See Algorithm 2 for more details.

StratPD and CatStratPD both have theoretical worst-case time complexity of $O(n^2)$ for n observations. For StratPD, stratification costs $O(pn\log n)$, computing y deltas for all observations among the leaves has linear cost, and averaging slopes across unique x_j ranges is on the order of $|unique(\mathbf{X}_j)| \times n$ or n^2 when all $\mathbf{X}_j$ are unique in the worst case. StratPD is, thus, $O(n^2)$ in the worst case. CatStratPD also stratifies in $O(pn\log n)$ and computes category deltas linearly in n but must make multiple passes over the $|T|$ leaves to average all possible leaf category delta vectors. In practice, three passes is the max we have seen (for high-cardinality variables), so we can assume the number of passes is some small

constant. Averaging two vectors costs $|unique(\mathbf{X}_j)|$, so each pass requires $|T| \times |unique(\mathbf{X}_j)|$. The number of leaves is roughly $n/\texttt{min\_samples\_leaf}$ and, worst-case, $|unique(\mathbf{X}_j)| = n$, meaning that merging dominates CATSTRATPD complexity leading to $O(n^2)$. Experiments show that our prototype is fast enough for practical use (see Section 4).

3 Related work

Friedman [2000] defines the partial dependence of $\hat{f}$ on feature subsets, S, as an expectation conditioned on the remaining variables:

$$FPD_S(x_S) = \mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})] = \frac{1}{n} \sum_{i=1}^{n} \hat{f}(x_S, x_{\backslash S}^{(i)})$$

The individual conditional expectation (ICE) plot from Goldstein et al. [2015] is a local method that estimates the partial dependence of the prediction $\hat{f}$ on x_S, or single variable x_j, across individual observations. Suppose that $(x_j^{(i)}, x_{\backslash j}^{(i)})$ are the values of the ith row of $\mathbf{X}$. For each i, the ICE plot produces a curve from the fitted model over all values of x_j while holding $x_{\backslash j}$ fixed: $\hat{f}_j^{(i)} = \hat{f}((\{x_j^{(k)}\}_{k=1}^n, x_{\backslash j}^{(i)}))$. By construction, the FPD curve for a variable x_j is the average over all ICE curves for that variable. The motivation for ICE is to identify variable dependencies not shown by the FPD curve (because they average out). Typically both ICE and FPD are used to examine the partial dependence of y on x_j. Both techniques operate by observing $\hat{f}$ output as $x_{\backslash j}$ (FPD) or x_j (ICE) values are shifted through all n observations.

Rather than shift feature values, SHAP from Lundberg and Lee [2017], with roots in *Shapley regression values* [Lipovetsky and Conklin, 2001], calculates the average marginal effect of adding x_j to models trained on all possible subsets of features, F:

$$\phi_j(\hat{f}, \mathbf{x}) = \sum_{S \subseteq F \backslash \{j\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} [\hat{f}_{S \cup \{j\}}(x_{S \cup \{j\}}) - \hat{f}_S(x_S)]$$

To avoid training a combinatorial explosion of models with the various feature subsets, $\hat{f}_S(x_S)$, SHAP reuses a single model fitted to $(\mathbf{X}, \mathbf{y})$ by running simplified feature vectors into the model. One possible implementation for simplified vectors is to replace "missing" x_j features with the expected value of x_j. SHAP uses a more general approach that approximates $\hat{f}_S(x_S)$ with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})|\mathbf{X}_S = x_S]$ where $\mathbf{X}$ is called the *background set* and users can pass in, for example, a single vector with x_j averages or even the entire training set. For efficiency reasons, SHAP's implementation (in "interventional" mode) approximates $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})|\mathbf{X}_S = x_S]$ with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\backslash S})]$. To further reduce computation costs, SHAP users typically explain a very small subsample of the data set, but with potentially a commensurate reduction in the explanatory resolution of the underlying population. SHAP has model-type-dependent optimizations for linear regression, deep learning, and decision-tree based models.

while the formal SHAP definition does not assume feature independence, the SHAP implementation does so for the specific purpose of simulating models trained on feature subsets. The SHAP definition approximates $\hat{f}_S(x_S)$ with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\bar{S}})|\mathbf{X}_S = x_S]$, but SHAP's implementation approximates $\hat{f}_S(x_S)$ with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\bar{S}})]$, which requires feature independence for accuracy. **?** argues that "*unconditional* [as implemented] rather than *conditional* [as defined] expectations provide the right notion of dropping features." In practice, however, we have observed several instances of questionable SHAP partial dependence curves for features important to the model.

From a mathematical perspective, the bias is clear because $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\bar{S}})]$ is equivalent to $FPD_S(x)$ that is known to be biased by codependent features. (Equation 53 in Friedman 2000 gives a summation equivalent to $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\bar{S}})]$)

ALE has same goal as us. we are model-free version. both estimate partial derivatives then cumsum. We do it point to point, they do in regions from min to max of support. both use finite differences; ALE using finite differences on $\hat{f}$ and us on $\mathbf{y}$. They do binning with intervals and we jump between unique x_j. that means they can ask for predictions for previously unseen x_j even if within the support range. By the way what if the support range is not a simple Gaussian density function? are we then picking values in low-density areas such as in between humps? anyway, we only consider existing values.

ALE cats appendix E: "The main consideration for a ?reasonable? ordering is that function extrapolation outside the data envelope should be avoided as much as possible when neighboring of values of xj are plugged into f(xj , x not j)." "we order the levels of Xj based on how dissimilar the sample values xi, not j : i = 1, 2, . . . , n are across the levels of Xj." ordering Boolean then has no effect since any order has the same difference between y values of true/false. K set to the number of nonempty category levels. from his email: "The way ALE plots treats categorical predictors still requires some extrapolation. With binary categorical variables, the extrapolation would be exactly the same as for PD plots. But with many categories, the extrapolation would be substantially less."

ALE guys says "the difference is that ALE plots hold $x_{\setminus j}$ exactly constant as it varies x_j across categories, which it can do because it uses the fitted model $f(x_j, x_{notj})$, whereas your approach keeps $x_{\setminus j}$ approximately the same as it varies x_j across categories, which it must do because it does not use any fitted model."


Figure 2: x_2 continuous and partitioned by decision tree fit to $(x_{\setminus j}, \mathbf{y})$, x_1 continuous (a), (b) and categorical (c), (d). Small wedges on x_1 axis indicate PD curve points

if x_2 is also cat then assuming same value means shifting x_2 to new cat as well for us, right? dec trees tend to test existing values for splits so will likely keep integers in their own leaves. In the case of x2 being cat with k levels, we can use `max_leaf_nodes=8, min_samples_leaf=1` but in general we can still shift x2 cats during strat. verified by experiment.

ALE partitions x_j into bins then fixes $x_{\setminus j}$ to compute average finite difference of y at $x_j = z$ in x_j neighborhood (all $x_{\setminus j}$ in that x_j bin) whereas we partition $x_{\setminus j}$ then compute average finite diff of y at $x_j = z$ across those $p - 1$ dim regions hoping they are repetitions of the same $p - 1$ dimensional vector (all same).

TODO: in case of discrete...

The inner difference $\hat{f}_{S \cup \{j\}}(x_{S \cup \{j\}}) - \hat{f}_S(x_S) = \mathbb{E}[\hat{f}(x_{S \cup \{j\}}, \mathbf{X}_{\setminus (S \cup \{j\})})] - \mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\setminus S})] = FPD_{S \cup \{j\}}(\mathbf{x}) - FPD_S(\mathbf{x})$

so in practice this keeps the same model but shifts feature values like FPD/ICE. SHAP also observes changes in model output for shifted feature values, but using models trained on all possible subsets, S, of features:

Comparison

$FPD_j(z) = \mathbb{E}[\hat{f}(x_j = z, \mathbf{X}_{\setminus j})])$

SHAP for one S is $\phi_j(\hat{f}, \mathbf{x}) = \mathbb{E}[\hat{f}(x_{S \cup \{j\}}, \mathbf{X}_{\setminus (S \cup \{j\})}) | \mathbf{X}_{S \cup \{j\}} = x_{S \cup \{j\}}] - \mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\setminus S}) | \mathbf{X}_S = x_S]$ for all subsets $S \subset F$ for $F = \{1, 2, .., p\}$. When $S = F \setminus \{j\}$, like ALE, it becomes $\mathbb{E}[\hat{f}(x_F)] - \mathbb{E}[\hat{f}(x_{F \setminus \{j\}}, \mathbf{X}_j)]$ or $\hat{f}(\mathbf{x}) - FPD_{F \setminus \{j\}}(x)$ if assume independence.

ALE at $x_j = z$, it is cumsum of $\mathbb{E}[\hat{f}(b_k, \mathbf{X}_{\setminus j}) - \hat{f}(b_{k-1}, \mathbf{X}_{\setminus j}) \,|\, x_j \in (b_{k-1}, b_k]]$ for bin b_k partitioning min..max for var j into K intervals.

StratPD at $x_j = z$ is $\mathbb{E}[(y^{(i_L+1)} - y^{i_L})/(x_j^{(i_L+1)} - x_j^{(i_L)}) \,|\, z \in [x_j^{(i_L+1)} - x_j^{(i_L)}) \text{ and } L \in T]$

$\mathbb{E}[\phi_j(\hat{f}, x)|\mathbf{X}_j = z] = FPD_j(x) - \bar{y}$ if features independent. SHAP's implementation approximates $\hat{f}$ trained on just S features, $\hat{f}_S(x_S)$, with $\mathbb{E}[\hat{f}(x_S, \mathbf{X}_{\bar{S}})]$ by assuming independent features.

Strobl thing Strobl et al. [2008]: "conditional permutation scheme, where X j is permuted only within groups of observations with Z = z, to preserve the correlation structure between X j and the other predictor variables." "Our suggestion is to define the grid within which the values of X j are permuted for each tree by means of the partition of the feature space induced by that tree. The main advantages of this approach are that this partition was already learned from the data during model fitting, contains splits in categorical, ordered and continuous predictor variables and can thus serve as an internally available means for discretizing the feature space."

Reviewer: "Their motivation is to handle the same extrapolation problems caused by having to plug in nonsensical combinations of correlated predictors, which is common to PD plots and permutation based importance measures."

cold start, counting execution time and number of hyper parameters. particularly deep learning

The techniques differ in algorithm simplicity, performance, and ability to isolate codependent variables. a nonparametric technique could also inform which machine learning model to use if a model is desired.

SHAP is mean centered FPD for independent variables, proof in supplemental material.

4 Experimental results

In this section, we demonstrate experimentally that StratPD and CatStratPD compute accurate partial dependence curves for synthetic data and plausible results for a real data set. Experiments also provide evidence that existing model-based techniques can provide meaningfully-biased curves. We begin by comparing the partial dependence curves from popular techniques on synthetic data with complex interactions.

Figure 3 illustrates that FPD, SHAP, ALE, and StratPD all suitably isolate the effect of independent individual variables on the response for noiseless data generated via: $y = x_1^2 + x_1 x_2 + 5x_1 sin(3x_2) + 10$ for $x_1, x_2, x_3 \sim U(0, 10)$ where x_3 does not affect y. The shapes of the curves for all techniques look similar except that StratPD starts all curves at $y = 0$ (as could the others). SHAP's curves have the advantage that they indicate the presence of variable interactions. To our eye, StratPD's curves are smoothest despite not having access to model predictions.


Figure 3: $y = x_1^2 + x_1 x_2 + 5x_1 sin(3x_2) + 10$ where $x_1, x_2, x_3 \sim U(0, 10)$ and x_3 does not affect y. No noise added. Random forest with 10 trees. R^2 0.9973

Models have a tendency to smooth out noise and a legitimate concern is that, without the benefit of a model, StratPD could be adversely affected. Figure 4 demonstrates StratPD

curves for noisy quadratics generated from $y = x_1^2 + x_1 + 10 + \epsilon$ where $\epsilon \sim N(0, \sigma)$ ϵ and, at $\sigma = 2$, 95% of the noise falls within [0,4] (since $2\sigma = 4$), meaning that the signal-to-noise ratio is at best 1-to-1 for x_1^2 in $[-2, 2]$. For zero-centered Gaussian noise and this data set, STRATPD accuracy is stable. The superfluous noise variable x_3 in Figure 3 also did not confuse STRATPD.


Figure 4: $y = x_1^2 + x_1 + 10 + \epsilon$ where $x_1, x_2 \sim U(-2, 2)$, $\epsilon \sim N(0, \sigma)$ with $\sigma \in [0, 0.5, 1, 2]$.

Turning to categorical variables, Figure 5 presents partial dependence curves for FPD, ALE, and CATSTRATPD derived from a noisy synthetic weather data set, where temperature varies in sinusoidal fashion over the year and with different baseline temperatures per state. (The vertical "smear" in the FPD plot shows the complete sine waves but from the side, edge on.) Variable `state` is independent and all plots identify the baseline temperature per state correctly.


Figure 5: $y = base[x_{state}] + 10 sin(\frac{2\pi}{365} x_{dayofyear} + \pi) + \epsilon$ where $\epsilon \sim N(0, 4)$. The baseline temperature per state is $\{AZ = 90, CA = 70, CO = 40, NV = 80, WA = 60\}$. Sinusoids in (a) are the average of three years' temperature data.

The primary goal of the stratification approach proposed in this paper is to obtain accurate partial dependence curves in the presence of codependent variables. To test STRATPD/CATSTRATPD and discover potential bias in existing techniques, we synthesized a body weight data set generated by the following equation with nontrivial codependencies between variables:

$$
\begin{aligned}
y &= 120 + 10(x_{height} - min(x_{height})) + 40 x_{pregnant} - 1.5 x_{education} \\
\text{where} \quad & x_{sex} \sim Bernoulli(\{M, F\}, p = 0.5) \\
& x_{pregnant} = \begin{cases} Bernoulli(\{0, 1\}, p = 0.5) & \text{if } x_{sex} = F \\ 0 & \text{if } x_{sex} = M \end{cases} \\
& x_{height} = \begin{cases} 5 * 12 + 5 + \epsilon & \text{if } x_{sex} = F, \ \epsilon \sim U(-4.5, 5) \\ 5 * 12 + 8 + \epsilon & \text{if } x_{sex} = M, \ \epsilon \sim U(-7, 8) \end{cases} \\
& x_{education} = \begin{cases} 12 + \epsilon & \text{if } x_{sex} = F, \ \epsilon \sim U(0, 8) \\ 10 + \epsilon & \text{if } x_{sex} = M, \ \epsilon \sim U(0, 8) \end{cases}
\end{aligned}
\tag{2}
$$

The partial derivative of y with respect to x_{height} is 10 (holding all other variables constant), so the optimal partial dependence curve is a line with slope 10. Figure 6 illustrates the curves for the techniques under consideration, with ALE and STRATPD giving the sharpest representation of the linear relationship. (STRATPD's curve is drawn on top of the SHAP plots using the righthand scale.) The FPD and both SHAP plots also suggest a linear relationship, albeit with a little less precision. The ICE curves in Figure 6(a) and "fuzzy" SHAP curves have the advantage that they alert users to variable dependencies or interaction terms. On the other hand, the kink in the partial dependence curve and other visual phenomena could confuse less experienced machine learning practitioners and certainly analysts and researchers in other fields (our primary target communities).

Figure 6: SHAP partial dependence plots of response body weight on feature `height` using 2000 synthetic observations from Equation 2. 50/50 male/female, half of the women are pregnant, and pregnancy contributes 40 pounds. SHAP interrogated an RF with 40 trees and explained 2000 samples; the interventional case used 100 observations as background data. RF OOB R^2 0.999


Figure 7: foo.

Even for experts, explaining this behavior requires some thought, and one must distinguish between model artifacts and interesting phenomena. The discontinuity at the maximum female height location arises partially from the model having trouble extrapolating for extremely tall pregnant women. Consider one of the upper ICE lines in Figure 6(a) for a pregnant woman. As the ICE line slides x_{height} above the maximum height for a woman, the model leaves the support of the training data and predicts a *lower* weight as height increases. ALE's curve is straight because it focuses on local effects, demonstrating that the lack of sharp slope-10 lines for FPD and SHAP cannot be attributed simply to a poor choice of model.

SHAP also defines x_j feature importance as the average magnitude of the x_j SHAP values, which introduces a paradox. The spread of the SHAP values alerts users to variable interactions, but allows contributions from other variables to leak in, thus, potentially leading to less precise estimates of x_{height}'s importance.

It is conceivable that a more sophisticated model (in terms of extrapolation) could sharpen the FPD and SHAP curves for x_{height}. There is a difference, however, between extrapolating to a meaningful but unsupported vector and making predictions for nonsensical vectors arising from variable codependencies. Techniques that rely on such predictions make the implicit assumption of variable independence, introducing the potential for bias. Consider Figure 7 that presents the partial dependence results for categorical variable $x_{pregnant}$ (same data set). The weight gain from pregnancy is 40 pounds per Equation 2, but only CATSTRATPD identifies that exact relationship; FPD, SHAP, and ALE show a gain of 30 pounds.

CATSTRATPD stratifies persons with the same or similar sex, education, and height into groups and then examines the relationship between $x_{pregnant}$ and y. If a group contains both pregnant and nonpregnant females, the difference in weight will be 40 pounds in this noiseless data set (assuming identical $x_{\setminus j}$). FPD and ALE rely on computations that require fitted models to conjure up predictions for nonsensical records, such as pregnant males (e.g., $\hat{f}(x_j = pregnant, \mathbf{X}_{\setminus j})$). Not even a human knows how to estimate the weight of a pregnant male. SHAP, per its definition, does not require such predictions, but in practice for efficiency reasons, SHAP approximates $\mathbb{E}[\hat{f}(x_j = pregnant, \mathbf{X}_{\setminus j})|\mathbf{X}_j = x_j]$ with $\mathbb{E}[\hat{f}(x_j = pregnant, \mathbf{X}_{\setminus j})]$, which does not restrict pregnancy to females. As discussed above, there are advantages to all of these model-based techniques, but this example demonstrates there is potential for partial dependence bias.

9

Figure 8: (a) Marginal plot of bulldozer `YearMade` versus `SalePrice` using subsample of 20k observations, (b) partial dependence drawn by SHAP interrogating an RF with 40 trees and explaining 1000 values with 100 observations as background data, (c) STRATPD partial dependence.

The stratification approach also gives plausible results for real data sets, such as the bulldozer auction data from Kaggle [2018]. Figure 8 shows the partial dependence curves for the same set of techniques as before on feature `YearMade`, chosen because it is very predictive of sale price. The shape and magnitude of the FPD, ALE, and STRATPD curves are similar, indicating that older bulldozers are worth less at auction, which is plausible. The STRATPD curve shows 10 bootstrapped trials where the heavy black dots represent the partial dependence curve and the other curves describe the variability. The SHAP curve looks unexpectedly random, but the underlying mathematics are sound; perhaps this identifies an issue arising from performance compromises. (We have submitted an issue to their repo).

And, finally, an important consideration for any tool is performance, so we plotted execution time versus data size (up to 30,000 observations) for three real Kaggle data sets: rent, bulldozer, and flight arrival delays Kaggle [2015]. Figure 9 shows growth curves for 40 numerical variables and 11 categorical variables grouped by type of variable. For these data sets, STRATPD takes 1.1s or less to process 30,000 records for any x_j, despite the potential for quadratic cost. CATSTRATPD typically processes categorical variables in less than 2s but takes 13s for the high-cardinality categorical `ModelID` of bulldozer (which looks mildly quadratic). These elapsed times for our prototype show it to be practical and competitive with FPD/ICE, SHAP, and ALE. If the cost to train a model using (cross validated) grid search for hyper parameter tuning is included, STRATPD and CATSTRATPD outperforms these existing techniques (as training and tuning is often measured in minutes).


Figure 9: timing blah. $p = 20$, bulldozer $p = 14$, and flight $p = 17$. Numba JIT warmup Numba just-in-time compiler (`numba.pydata.org`)

what if X,y relationship is very weak? models would get low accuracy. what happens to us? I think we would simply show low partial dependence curves.

TODO: unsupervised?

All simulations in this section were run on a 4.0 Ghz 32G RAM machine running OS X 10.13.6 with SHAP 0.34, scikit-learn 0.21.3, XGBoost 0.90, and Python 3.7.4.

5   Discussion and future work

can parallelize code or do in C++.

similar to ALE but much simpler for categorical variables.

we can handle partial derivatives for $x_{j,l}$ using 2D finite difference to get gradients then using double integral instead of single cumulative sum. I believe that gets us an estimate for plot value at point (x_j, x_l). actually looks like the ALE guy is calling it a second order effect and is subtracting the two one-dimensional finite differences. Hmm... he's just holding x_j, x_l constant and integrating over the other variables to find the value at (x_j, x_l).

it couldn't handle it when I got rid of education in body weight data set. got 0 pdpy. perhaps investigate

A word on hyper parameters. All techniques have hyper parameters that can affect results, particularly if one includes hyper parameters in the model itself for model-based techniques. For example, a SHAP plot like Figure 8(c) with fewer explained vectors looks much more like a line but then does not alert the user about variable interactions. It is sound practice to try a variety of hyper parameters with any partial dependence technique, including ours.

Our goal here is not to say these other techniques are not useful. Rather, we are pointing out potential issues and hoping to open a new line of inquiry that experiments have shown to be accurate in cases where previous techniques are not.

6 Appendix

Algorithm 1: $StratPD(\mathbf{X}, \mathbf{y}, j, min\_samples\_leaf, min\_slopes\_per\_x)$

$T :=$ Decision tree regressor fit to $(x_{\setminus j}, \mathbf{y})$ with hyper parameter: $min\_samples\_leaf$
for *each leaf* $L \in T$ **do**
  $(\mathbf{x}_L, \mathbf{y}_L) := \{(x_j^{(i)}, y^{(i)})\}_{i \in L}$ ▷ *Get leaf observations*
  $\mathbf{ux}, \bar{\mathbf{y}} :=$ Group and sort $(\mathbf{x}_L, \mathbf{y}_L)$ by x_j value, computing $\bar{y}$ per unique x_j value
  $\mathbf{dx} := \mathbf{ux}^{(i+1)} - \mathbf{ux}^{(i)}$ for $i = 1..|\mathbf{ux}| - 1$ ▷ *Discrete difference between adjacent*
  $\mathbf{dy} := \bar{\mathbf{y}}^{(i+1)} - \bar{\mathbf{y}}^{(i)}$ for $i = 1..|\mathbf{ux}| - 1$
  Add tuples $(\mathbf{ux}^{(i)}, \mathbf{ux}^{(i+1)}, \mathbf{dy}^{(i)}/\mathbf{dx}^{(i)})$ to list $\mathbf{D}$ for $i = 1..|\mathbf{ux}| - 1$
$\mathbf{ux} := unique(\mathbf{X}_j)$
for *each* $x \in \mathbf{ux}$ **do** ▷ *Count slopes and compute average slope per unique x_j value*
  $slopes := [slope$ for $(a, b, slope) \in \mathbf{D}$ if $x \geq a$ and $x < b]$
  $\mathbf{c}_x := |\mathbf{slopes}|$
  $\mathbf{dydx}_x := \overline{\mathbf{slopes}}$
$\mathbf{dydx} := \mathbf{dydx}[\mathbf{c} \geq min\_slopes\_per\_x]$ ▷ *Drop missing slopes, those computed with too few*
$\mathbf{ux} := \mathbf{ux}[\mathbf{c} \geq min\_slopes\_per\_x]$
$\mathbf{pdx} := \mathbf{ux}^{(i+1)} - \mathbf{ux}^{(i)}$ for $i = 1..|\mathbf{ux}| - 1$
$\mathbf{pdy} := [0] + cumulative\_sum(\mathbf{dydx} * \mathbf{pdx})$ ▷ *integrate, inserting 0 for leftmost x_j*
return $\mathbf{pdx}, \mathbf{pdy}$

Algorithm 2: $CatStratPD(\mathbf{X}, \mathbf{y}, j, min\_samples\_leaf)$

$T :=$ Decision tree regressor fit to $(x_{\setminus j}, \mathbf{y})$ with hyper parameter: $min\_samples\_leaf$
Let ΔY be a matrix whose columns are vectors of leaf category deltas
Let C be a matrix whose columns are vectors for leaf category counts
for *each leaf* $L \in T$ **do** ▷ *Get average y delta relative to random ref category for obs. in leaves*
  $(\mathbf{x}_L, \mathbf{y}_L) := \{(x_j^{(i)}, y^{(i)})\}_{i \in L}$ ▷ *Get leaf observations*
  $ucats, C_L := unique(\mathbf{x}_L)$ ▷ *Get unique categories, counts from leaf observations*
  $\bar{\mathbf{y}} :=$ Group leaf records $(\mathbf{x}_L, \mathbf{y}_L)$ by category, computing $\bar{y}$ per unique category
  $refcat :=$ random category from $\mathbf{x}_L$
  $\Delta Y_L = \bar{\mathbf{y}} - \bar{\mathbf{y}}_{refcat}$
$\Delta \mathbf{y}, \mathbf{c} := \Delta Y_1, C_1$ ▷ *$\Delta \mathbf{y}$ is running average vector mapping category to average y delta*
$completed := \{1\}; work := \{2..|leaves|\};$ ▷ *set of leaf indexes*
while $|work| > 0$ and $|completed| > 0$ **do** ▷ *2 passes is typical to merge all ΔY_L into $\Delta \mathbf{y}$*
  $completed := \emptyset$
  for *each leaf* L *in* $work$ **do**
    if ΔY_L *has category in common with* $\Delta \mathbf{y}$ **then**
      $completed := completed \cup \{L\}$
      $c =$ random category in intersection
      $\Delta \mathbf{y}_L := \Delta Y_L - \Delta Y_{L,c} + \Delta \mathbf{y}_c$ ▷ *Adjust so $\Delta Y_{L,c} = 0$, add corresponding $\Delta \mathbf{y}_c$ value*
      $\Delta \mathbf{y} := (\mathbf{c} \times \Delta \mathbf{y} + C_L \times \Delta \mathbf{y}_L)/(\mathbf{c} + C_L)$ where $z + NaN = z$ ▷ *weighted average*
      $\mathbf{c} := \mathbf{c} + C_L$ ▷ *update running weight*
  $work := work \setminus completed$
return $\Delta \mathbf{y}$

References

Jerome H. Friedman. Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29:1189–1232, 2000.

Giles Hooker. Generalized functional anova diagnostics for high-dimensional functions of dependent variables. *Journal of Computational and Graphical Statistics*, 16(3):709–732, 2007. doi: 10.1198/106186007X237892. URL `https://doi.org/10.1198/106186007X237892`.

Alex Goldstein, Adam Kapelner, Justin Bleich, and Emil Pitkin. Peeking inside the black box: Visualizing statistical learning with plots of individual conditional expectation. *Journal of Computational and Graphical Statistics*, 24(1):44–65, 2015. doi: 10.1080/10618600.2014.907095. URL `https://doi.org/10.1080/10618600.2014.907095`.

Daniel W Apley. Visualizing the effects of predictor variables in black box supervised learning models. *arXiv preprint arXiv:1612.08468*, 2016.

Scott M Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems 30*, pages 4765–4774. Curran Associates, Inc., 2017. URL `http://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions.pdf`.

Kaggle. Two sigma connect: Rental listing inquiries. `https://www.kaggle.com/c/two-sigma-connect-rental-listing-inquiries`, 2017. Accessed: 2019-06-15.

Stan Lipovetsky and Michael Conklin. Analysis of regression in game theory approach. *Applied Stochastic Models in Business and Industry*, 17:319 − 330, 10 2001. doi: 10.1002/asmb.446.

Carolin Strobl, Anne-Laure Boulesteix, Thomas Kneib, Thomas Augustin, and Achim Zeileis. Conditional variable importance for random forests. *BMC Bioinformatics*, 9(1):307, 2008. doi: 10.1186/1471-2105-9-307. URL `https://doi.org/10.1186/1471-2105-9-307`.

Kaggle. Blue book for bulldozer. `https://www.kaggle.com/sureshsubramaniam/blue-book-for-bulldozer-kaggle-competition`, 2018. Accessed: 2019-06-15.

Kaggle. 2015 flight delays and cancellations. `https://www.kaggle.com/usdot/flight-delays`, 2015. Accessed: 2019-12-01.